

dog_app-jm

January 27, 2020

1 Convolutional Neural Networks

1.1 Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets: * Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dogImages.

- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays `human_files` and `dog_files`.

```
In [23]: import numpy as np
         from glob import glob

         # load filenames for human and dog images
         human_files = np.array(glob("/content/drive/My Drive/Colab Notebooks/data/lfw/**/*.jpg"))
         dog_files = np.array(glob("/content/drive/My Drive/Colab Notebooks/data/dogImages/**/*.jpg"))

         # print number of images in each dataset
         print('There are %d total human images.' % len(human_files))
         print('There are %d total dog images.' % len(dog_files))
```

There are 7653 total human images.

There are 8351 total dog images.

Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the `haarcascades` directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [24]: import cv2
         import matplotlib.pyplot as plt
         %matplotlib inline

         # extract pre-trained face detector
         face_cascade = cv2.CascadeClassifier('/content/drive/My Drive/Colab Notebooks/data/haarcascades/haarcascade_frontalface_default.xml')

         # load color (BGR) image
         img = cv2.imread(human_files[0])
         # convert BGR image to grayscale
         gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

         # find faces in image
         faces = face_cascade.detectMultiScale(gray)

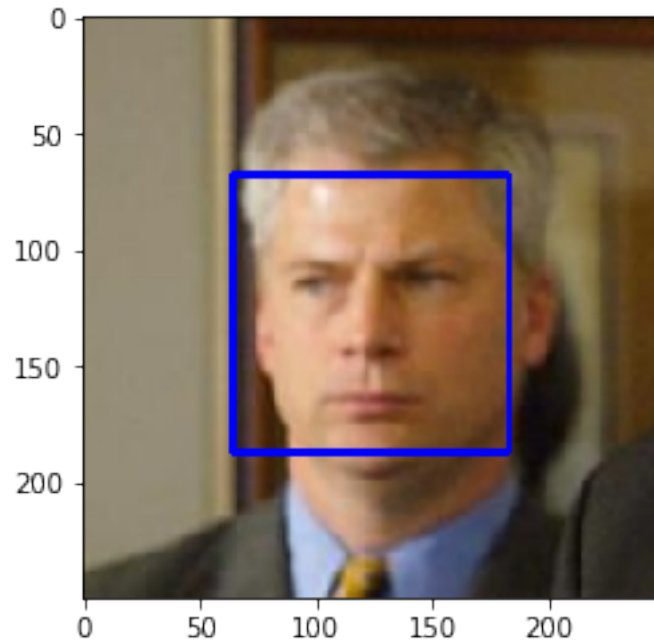
         # print number of faces detected in the image
         print('Number of faces detected:', len(faces))

         # get bounding box for each detected face
         for (x,y,w,h) in faces:
             # add bounding box to color image
             cv2.rectangle(img, (x,y), (x+w,y+h), (255,0,0), 2)
```

```
# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()
```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns `True` if a human face is detected in an image and `False` otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [0]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
```

```

img = cv2.imread(img_path)
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
faces = face_cascade.detectMultiScale(gray)
return len(faces) > 0

```

1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer: The algorithm detected human face for 98% images of humans and 7% images of dogs.

```

In [0]: from tqdm import tqdm
import collections

human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

##-## Do NOT modify the code above this line. ##-##

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.

counter = collections.Counter()

for path in human_files_short:
    if face_detector(path):
        counter["human"] += 1

for path in dog_files_short:
    if face_detector(path):
        counter["dog"] += 1

print("The algorithm detected human face for {0}% images of humans and {1}% images of dogs.")

```

The algorithm detected human face for 98% images of humans and 7% images of dogs.

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [0]: ### (Optional)  
### TODO: Test performance of another face detection algorithm.  
### Feel free to use as many code cells as needed.
```

Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of 1000 categories.

```
In [0]: import torch  
import torchvision.models as models  
  
# define VGG16 model  
VGG16 = models.vgg16(pretrained=True)  
  
# check if CUDA is available  
use_cuda = torch.cuda.is_available()  
  
# move model to GPU if CUDA is available  
if use_cuda:  
    VGG16 = VGG16.cuda()
```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as 'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg') as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```
In [0]: from PIL import Image  
import torchvision.transforms as transforms  
from torch.autograd import Variable  
  
# Set PIL to be tolerant of image files that are truncated.  
from PIL import ImageFile  
ImageFile.LOAD_TRUNCATED_IMAGES = True
```

```

def VGG16_predict(img_path):
    """
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path

    Args:
        img_path: path to an image

    Returns:
        Index corresponding to VGG-16 model's prediction
    """

    ## TODO: Complete the function.
    ## Load and pre-process an image from the given img_path
    ## Return the *index* of the predicted class for that image

    image = Image.open(img_path)
    transform = transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    ])

    image = transform(image)
    image = image.unsqueeze(0)
    if use_cuda:
        image = image.cuda()
    prediction = VGG16(image)

    return np.argmax(prediction.cpu().detach().numpy()) # predicted class index

```

1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the `dog_detector` function below, which returns True if a dog is detected in an image (and False if not).

```

In [0]: """ returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    ## TODO: Complete the function.

    result = VGG16_predict(img_path)

    return result >= 151 and result <= 268

```

1.1.6 (IMPLEMENTATION) Assess the Dog Detector

Question 2: Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

Answer: The algorithm detected dog for 0% images of humans and 97% images of dogs.

```
In [0]: ### TODO: Test the performance of the dog_detector function  
### on the images in human_files_short and dog_files_short.  
  
counter = collections.Counter()  
  
for path in human_files_short:  
    if dog_detector(path):  
        counter["human"] += 1  
  
for path in dog_files_short:  
    if dog_detector(path):  
        counter["dog"] += 1  
  
print("The algorithm detected dog for {0}% images of humans and {1}% images of dogs.".format(counter["human"], counter["dog"]))
```

The algorithm detected dog for 0% images of humans and 97% images of dogs.

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [0]: ### (Optional)  
### TODO: Report the performance of another pre-trained network.  
### Feel free to use as many code cells as needed.
```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dogImages/train`, `dogImages/valid`, and `dogImages/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [0]: import os
        from torchvision import datasets
        from torchvision import transforms
        from torch.utils import data

        import matplotlib.pyplot as plt

        ### TODO: Write data loaders for training, validation, and test sets
        ## Specify appropriate transforms, and batch_sizes

        data_transform = transforms.Compose([
            transforms.RandomResizedCrop(224),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
        ])

        data_transform_train = transforms.Compose([
            transforms.RandomResizedCrop(224),
            transforms.RandomRotation((0, 90)),
            transforms.ToTensor(),
```



```

        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    ])

    loaders_scratch = {}
    batch_size = 10
    num_workers = 2

    train_dataset = datasets.ImageFolder(root='/content/drive/My Drive/Colab Notebooks/data/')
    loaders_scratch['train'] = data.DataLoader(train_dataset, batch_size=batch_size, shuffle=True)

    valid_dataset = datasets.ImageFolder(root='/content/drive/My Drive/Colab Notebooks/data/')
    loaders_scratch['valid'] = data.DataLoader(valid_dataset, batch_size=batch_size, shuffle=True)

    test_dataset = datasets.ImageFolder(root='/content/drive/My Drive/Colab Notebooks/data/')
    loaders_scratch['test'] = data.DataLoader(test_dataset, batch_size=batch_size, shuffle=True)

```

Question 3: Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

Answer:

Each image was resized and randomly cropped to a square of 224x224 pixels. Following that the image was normalized according to library instructions. Such a size was selected because data loss could not be seen and image was not too big. I decided to augment training data because after few tries to learn the network detection accuracy was not so great. After that I could see improvement in learning quality. As data augmentation I selected random rotation up to 90 degrees.

1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```

In [0]: import torch.nn as nn
import torch.nn.functional as F

# define the CNN architecture
class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN

        self.layer1 = nn.Sequential(
            nn.Conv2d(3, 16, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2))

        self.layer2 = nn.Sequential(

```

```

        nn.Conv2d(16, 32, kernel_size=3, stride=1, padding=1),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2, stride=2))

    self.layer3 = nn.Sequential(
        nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2, stride=2))

    self.layer4 = nn.Sequential(
        nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2, stride=2))

    self.dropout2 = nn.Dropout(0.2)
    self.dropout4 = nn.Dropout(0.4)

    self.full1 = nn.Sequential(
        nn.Linear(128 * 14 * 14, 512),
        nn.ReLU())

    self.fc2 = nn.Linear(512, 133)

def forward(self, x):
    ## Define forward behavior

    x = self.layer1(x)
    x = self.dropout2(x)
    x = self.layer2(x)
    x = self.dropout2(x)
    x = self.layer3(x)
    x = self.dropout2(x)
    x = self.layer4(x)

    x = x.view(-1, 128 * 14 * 14)

    x = self.dropout4(x)
    x = self.full1(x)
    x = self.dropout4(x)
    x = self.fc2(x)

    return x

### You do NOT have to modify the code below this line. ###

# instantiate the CNN
model_scratch = Net()

```

```
# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()
```

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer: In feature part of the network I decided to use 4 layers which consisted of: - convolutional layer with kernel size 3 increasing number of features - ReLu activation function - pooling layer to reduce image size by 2 - dropout mechanism with probability 0.2

In fully connected part I decided to use 2 layers. First reduced number of inputs from 25088 to 512 and the second one reduced number of inputs from 512 to number of classes to be detected - 133. There was dropout layer in between with probability of 0.4.

Such a model was selected because similar approach was used during the classes to similar problem. In addition, some of the network models used in pytorch were similar although usually more complicated. I tried to select simple model due to small amount of data.

1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```
In [0]: import torch.optim as optim

learning_rate = 0.001

### TODO: select loss function
criterion_scratch = nn.CrossEntropyLoss()

### TODO: select optimizer
optimizer_scratch = optim.SGD(model_scratch.parameters(), lr=learning_rate)
```

1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_scratch.pt'`.

```
In [0]: # the following import is required for training to be robust to truncated images
import numpy as np
from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True

def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf
    correct = 0.
    total = 0.

    for epoch in range(1, n_epochs+1):
```

```

# initialize variables to monitor training and validation loss
train_loss = 0.0
valid_loss = 0.0

#####
# train the model #
#####
model.train()
for batch_idx, (data, target) in enumerate(loaders['train']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    ## find the loss and update the model parameters accordingly
    ## record the average training loss, using something like
    ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

    # clear the gradients of all optimized variables
    optimizer.zero_grad()
    # forward pass: compute predicted outputs by passing inputs to the model
    output = model(data)
    # calculate the batch loss
    loss = criterion(output, target)
    # backward pass: compute gradient of the loss with respect to model parameters
    loss.backward()
    # perform a single optimization step (parameter update)
    optimizer.step()
    # update training loss
    train_loss += ((1 / (batch_idx + 1)) * (loss.data - train_loss))

#####
# validate the model #
#####
model.eval()
for batch_idx, (data, target) in enumerate(loaders['valid']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    ## update the average validation loss

    # forward pass: compute predicted outputs by passing inputs to the model
    output = model(data)
    # calculate the batch loss
    loss = criterion(output, target)
    # update average validation loss
    valid_loss += ((1 / (batch_idx + 1)) * (loss.data - valid_loss))
    # convert output probabilities to predicted class
    pred = output.data.max(1, keepdim=True)[1]

```

```

        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy)
        total += data.size(0)

    # print training/validation statistics
    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f} \tValidation
          epoch,
          train_loss,
          valid_loss,
          100. * correct / total
    ))

    ## TODO: save the model if validation loss has decreased
    if valid_loss < valid_loss_min:
        torch.save(model.state_dict(), save_path)
        valid_loss_min = valid_loss

    # return trained model
    return model

# train the model
model_scratch = train(400, loaders_scratch, model_scratch, optimizer_scratch,
                      criterion_scratch, use_cuda, '/content/drive/My Drive/Colab Notebo

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('/content/drive/My Drive/Colab Notebooks/model_

```

Epoch: 1	Training Loss: 4.889802	Validation Loss: 4.887720	Validation Ac
Epoch: 2	Training Loss: 4.880711	Validation Loss: 4.884031	Validation Ac
Epoch: 3	Training Loss: 4.874452	Validation Loss: 4.880095	Validation Ac
Epoch: 4	Training Loss: 4.872150	Validation Loss: 4.878139	Validation Ac
Epoch: 5	Training Loss: 4.866072	Validation Loss: 4.874239	Validation Ac
Epoch: 6	Training Loss: 4.861391	Validation Loss: 4.871423	Validation Ac
Epoch: 7	Training Loss: 4.860166	Validation Loss: 4.868813	Validation Ac
Epoch: 8	Training Loss: 4.853275	Validation Loss: 4.864301	Validation Ac
Epoch: 9	Training Loss: 4.844520	Validation Loss: 4.859428	Validation Ac
Epoch: 10	Training Loss: 4.831635	Validation Loss: 4.853657	Validation A
Epoch: 11	Training Loss: 4.819396	Validation Loss: 4.842750	Validation A
Epoch: 12	Training Loss: 4.812555	Validation Loss: 4.830974	Validation A
Epoch: 13	Training Loss: 4.792469	Validation Loss: 4.818999	Validation A
Epoch: 14	Training Loss: 4.779333	Validation Loss: 4.812827	Validation A
Epoch: 15	Training Loss: 4.778947	Validation Loss: 4.810057	Validation A
Epoch: 16	Training Loss: 4.770101	Validation Loss: 4.802031	Validation A
Epoch: 17	Training Loss: 4.765337	Validation Loss: 4.789282	Validation A
Epoch: 18	Training Loss: 4.761455	Validation Loss: 4.792146	Validation A
Epoch: 19	Training Loss: 4.737973	Validation Loss: 4.785418	Validation A

Epoch: 20	Training Loss: 4.739319	Validation Loss: 4.789217	Validation A
Epoch: 21	Training Loss: 4.726267	Validation Loss: 4.771544	Validation A
Epoch: 22	Training Loss: 4.737859	Validation Loss: 4.789311	Validation A
Epoch: 23	Training Loss: 4.719619	Validation Loss: 4.772806	Validation A
Epoch: 24	Training Loss: 4.714355	Validation Loss: 4.764539	Validation A
Epoch: 25	Training Loss: 4.707738	Validation Loss: 4.769601	Validation A
Epoch: 26	Training Loss: 4.702507	Validation Loss: 4.759590	Validation A
Epoch: 27	Training Loss: 4.686752	Validation Loss: 4.754711	Validation A
Epoch: 28	Training Loss: 4.681721	Validation Loss: 4.736724	Validation A
Epoch: 29	Training Loss: 4.665602	Validation Loss: 4.728213	Validation A
Epoch: 30	Training Loss: 4.651689	Validation Loss: 4.731114	Validation A
Epoch: 31	Training Loss: 4.639747	Validation Loss: 4.686635	Validation A
Epoch: 32	Training Loss: 4.614209	Validation Loss: 4.687821	Validation A
Epoch: 33	Training Loss: 4.612519	Validation Loss: 4.683485	Validation A
Epoch: 34	Training Loss: 4.588634	Validation Loss: 4.659617	Validation A
Epoch: 35	Training Loss: 4.589596	Validation Loss: 4.668017	Validation A
Epoch: 36	Training Loss: 4.579241	Validation Loss: 4.631346	Validation A
Epoch: 37	Training Loss: 4.573422	Validation Loss: 4.655549	Validation A
Epoch: 38	Training Loss: 4.551734	Validation Loss: 4.619094	Validation A
Epoch: 39	Training Loss: 4.558963	Validation Loss: 4.625632	Validation A
Epoch: 40	Training Loss: 4.529906	Validation Loss: 4.597682	Validation A
Epoch: 41	Training Loss: 4.534218	Validation Loss: 4.611000	Validation A
Epoch: 42	Training Loss: 4.524575	Validation Loss: 4.585309	Validation A
Epoch: 43	Training Loss: 4.523911	Validation Loss: 4.624469	Validation A
Epoch: 44	Training Loss: 4.513161	Validation Loss: 4.608538	Validation A
Epoch: 45	Training Loss: 4.510570	Validation Loss: 4.579090	Validation A
Epoch: 46	Training Loss: 4.492032	Validation Loss: 4.594042	Validation A
Epoch: 47	Training Loss: 4.486064	Validation Loss: 4.586446	Validation A
Epoch: 48	Training Loss: 4.479362	Validation Loss: 4.578300	Validation A
Epoch: 49	Training Loss: 4.474050	Validation Loss: 4.568787	Validation A
Epoch: 50	Training Loss: 4.477848	Validation Loss: 4.564169	Validation A
Epoch: 51	Training Loss: 4.461338	Validation Loss: 4.576610	Validation A
Epoch: 52	Training Loss: 4.446775	Validation Loss: 4.553746	Validation A
Epoch: 53	Training Loss: 4.451097	Validation Loss: 4.539131	Validation A
Epoch: 54	Training Loss: 4.445972	Validation Loss: 4.542055	Validation A
Epoch: 55	Training Loss: 4.439577	Validation Loss: 4.543793	Validation A
Epoch: 56	Training Loss: 4.432490	Validation Loss: 4.550891	Validation A
Epoch: 57	Training Loss: 4.423673	Validation Loss: 4.563600	Validation A
Epoch: 58	Training Loss: 4.411632	Validation Loss: 4.514838	Validation A
Epoch: 59	Training Loss: 4.408273	Validation Loss: 4.528983	Validation A
Epoch: 60	Training Loss: 4.414958	Validation Loss: 4.522494	Validation A
Epoch: 61	Training Loss: 4.376550	Validation Loss: 4.531257	Validation A
Epoch: 62	Training Loss: 4.381496	Validation Loss: 4.524548	Validation A
Epoch: 63	Training Loss: 4.383953	Validation Loss: 4.505863	Validation A
Epoch: 64	Training Loss: 4.377995	Validation Loss: 4.511341	Validation A
Epoch: 65	Training Loss: 4.372977	Validation Loss: 4.516416	Validation A
Epoch: 66	Training Loss: 4.371068	Validation Loss: 4.496193	Validation A
Epoch: 67	Training Loss: 4.357546	Validation Loss: 4.477891	Validation A

Epoch: 68	Training Loss: 4.359361	Validation Loss: 4.514221	Validation A
Epoch: 69	Training Loss: 4.358708	Validation Loss: 4.481921	Validation A
Epoch: 70	Training Loss: 4.342451	Validation Loss: 4.494312	Validation A
Epoch: 71	Training Loss: 4.336425	Validation Loss: 4.471353	Validation A
Epoch: 72	Training Loss: 4.337156	Validation Loss: 4.462639	Validation A
Epoch: 73	Training Loss: 4.331333	Validation Loss: 4.463763	Validation A
Epoch: 74	Training Loss: 4.324495	Validation Loss: 4.511412	Validation A
Epoch: 75	Training Loss: 4.330863	Validation Loss: 4.460423	Validation A
Epoch: 76	Training Loss: 4.310118	Validation Loss: 4.447963	Validation A
Epoch: 77	Training Loss: 4.295402	Validation Loss: 4.438911	Validation A
Epoch: 78	Training Loss: 4.304292	Validation Loss: 4.478023	Validation A
Epoch: 79	Training Loss: 4.298469	Validation Loss: 4.451281	Validation A
Epoch: 80	Training Loss: 4.302401	Validation Loss: 4.428257	Validation A
Epoch: 81	Training Loss: 4.288504	Validation Loss: 4.407689	Validation A
Epoch: 82	Training Loss: 4.292784	Validation Loss: 4.445201	Validation A
Epoch: 83	Training Loss: 4.284618	Validation Loss: 4.402153	Validation A
Epoch: 84	Training Loss: 4.265228	Validation Loss: 4.439430	Validation A
Epoch: 85	Training Loss: 4.280429	Validation Loss: 4.411247	Validation A
Epoch: 86	Training Loss: 4.265531	Validation Loss: 4.446830	Validation A
Epoch: 87	Training Loss: 4.277637	Validation Loss: 4.453113	Validation A
Epoch: 88	Training Loss: 4.246767	Validation Loss: 4.446455	Validation A
Epoch: 89	Training Loss: 4.250443	Validation Loss: 4.413422	Validation A
Epoch: 90	Training Loss: 4.256778	Validation Loss: 4.449054	Validation A
Epoch: 91	Training Loss: 4.248568	Validation Loss: 4.411288	Validation A
Epoch: 92	Training Loss: 4.236390	Validation Loss: 4.399438	Validation A
Epoch: 93	Training Loss: 4.231069	Validation Loss: 4.380850	Validation A
Epoch: 94	Training Loss: 4.226260	Validation Loss: 4.399861	Validation A
Epoch: 95	Training Loss: 4.241671	Validation Loss: 4.394013	Validation A
Epoch: 96	Training Loss: 4.235939	Validation Loss: 4.401415	Validation A
Epoch: 97	Training Loss: 4.207224	Validation Loss: 4.398232	Validation A
Epoch: 98	Training Loss: 4.210842	Validation Loss: 4.397831	Validation A
Epoch: 99	Training Loss: 4.203180	Validation Loss: 4.404888	Validation A
Epoch: 100	Training Loss: 4.204020	Validation Loss: 4.367791	Validation
Epoch: 101	Training Loss: 4.190968	Validation Loss: 4.377901	Validation
Epoch: 102	Training Loss: 4.181694	Validation Loss: 4.345314	Validation
Epoch: 103	Training Loss: 4.192420	Validation Loss: 4.383793	Validation
Epoch: 104	Training Loss: 4.183096	Validation Loss: 4.361443	Validation
Epoch: 105	Training Loss: 4.175605	Validation Loss: 4.361559	Validation
Epoch: 106	Training Loss: 4.178320	Validation Loss: 4.367634	Validation
Epoch: 107	Training Loss: 4.159990	Validation Loss: 4.353925	Validation
Epoch: 108	Training Loss: 4.147770	Validation Loss: 4.333640	Validation
Epoch: 109	Training Loss: 4.151306	Validation Loss: 4.317529	Validation
Epoch: 110	Training Loss: 4.148232	Validation Loss: 4.317950	Validation
Epoch: 111	Training Loss: 4.143281	Validation Loss: 4.364029	Validation
Epoch: 112	Training Loss: 4.126160	Validation Loss: 4.356215	Validation
Epoch: 113	Training Loss: 4.128595	Validation Loss: 4.335250	Validation
Epoch: 114	Training Loss: 4.115768	Validation Loss: 4.343791	Validation
Epoch: 115	Training Loss: 4.134300	Validation Loss: 4.289955	Validation

Epoch: 116	Training Loss: 4.124463	Validation Loss: 4.347414	Validation
Epoch: 117	Training Loss: 4.108920	Validation Loss: 4.339469	Validation
Epoch: 118	Training Loss: 4.091233	Validation Loss: 4.308765	Validation
Epoch: 119	Training Loss: 4.096913	Validation Loss: 4.282137	Validation
Epoch: 120	Training Loss: 4.072002	Validation Loss: 4.277309	Validation
Epoch: 121	Training Loss: 4.089199	Validation Loss: 4.335621	Validation
Epoch: 122	Training Loss: 4.070497	Validation Loss: 4.321894	Validation
Epoch: 123	Training Loss: 4.079008	Validation Loss: 4.284268	Validation
Epoch: 124	Training Loss: 4.052811	Validation Loss: 4.293931	Validation
Epoch: 125	Training Loss: 4.044912	Validation Loss: 4.251979	Validation
Epoch: 126	Training Loss: 4.063595	Validation Loss: 4.272233	Validation
Epoch: 127	Training Loss: 4.043044	Validation Loss: 4.296424	Validation
Epoch: 128	Training Loss: 4.037798	Validation Loss: 4.264029	Validation
Epoch: 129	Training Loss: 4.029923	Validation Loss: 4.291706	Validation
Epoch: 130	Training Loss: 4.035285	Validation Loss: 4.245619	Validation
Epoch: 131	Training Loss: 4.039813	Validation Loss: 4.212646	Validation
Epoch: 132	Training Loss: 4.009869	Validation Loss: 4.274029	Validation
Epoch: 133	Training Loss: 4.025888	Validation Loss: 4.260013	Validation
Epoch: 134	Training Loss: 4.016978	Validation Loss: 4.254373	Validation
Epoch: 135	Training Loss: 4.021310	Validation Loss: 4.253362	Validation
Epoch: 136	Training Loss: 4.003283	Validation Loss: 4.215093	Validation
Epoch: 137	Training Loss: 3.988268	Validation Loss: 4.256962	Validation
Epoch: 138	Training Loss: 3.976783	Validation Loss: 4.275323	Validation
Epoch: 139	Training Loss: 3.991629	Validation Loss: 4.244455	Validation
Epoch: 140	Training Loss: 3.977470	Validation Loss: 4.188062	Validation
Epoch: 141	Training Loss: 3.964296	Validation Loss: 4.218343	Validation
Epoch: 142	Training Loss: 3.949083	Validation Loss: 4.191419	Validation
Epoch: 143	Training Loss: 3.958943	Validation Loss: 4.218114	Validation
Epoch: 144	Training Loss: 3.951160	Validation Loss: 4.188188	Validation
Epoch: 145	Training Loss: 3.944432	Validation Loss: 4.183220	Validation
Epoch: 146	Training Loss: 3.936501	Validation Loss: 4.212515	Validation
Epoch: 147	Training Loss: 3.925274	Validation Loss: 4.181818	Validation
Epoch: 148	Training Loss: 3.931989	Validation Loss: 4.235027	Validation
Epoch: 149	Training Loss: 3.923272	Validation Loss: 4.209925	Validation
Epoch: 150	Training Loss: 3.911864	Validation Loss: 4.197152	Validation
Epoch: 151	Training Loss: 3.923867	Validation Loss: 4.176208	Validation
Epoch: 152	Training Loss: 3.914864	Validation Loss: 4.196151	Validation
Epoch: 153	Training Loss: 3.901565	Validation Loss: 4.166702	Validation
Epoch: 154	Training Loss: 3.898861	Validation Loss: 4.163126	Validation
Epoch: 155	Training Loss: 3.897008	Validation Loss: 4.158103	Validation
Epoch: 156	Training Loss: 3.882429	Validation Loss: 4.182452	Validation
Epoch: 157	Training Loss: 3.881769	Validation Loss: 4.127035	Validation
Epoch: 158	Training Loss: 3.856653	Validation Loss: 4.127237	Validation
Epoch: 159	Training Loss: 3.864453	Validation Loss: 4.138508	Validation
Epoch: 160	Training Loss: 3.857637	Validation Loss: 4.148266	Validation
Epoch: 161	Training Loss: 3.857973	Validation Loss: 4.145984	Validation
Epoch: 162	Training Loss: 3.830939	Validation Loss: 4.138818	Validation
Epoch: 163	Training Loss: 3.855467	Validation Loss: 4.147429	Validation

Epoch: 164	Training Loss: 3.837106	Validation Loss: 4.156875	Validation
Epoch: 165	Training Loss: 3.831790	Validation Loss: 4.119478	Validation
Epoch: 166	Training Loss: 3.820386	Validation Loss: 4.124229	Validation
Epoch: 167	Training Loss: 3.804905	Validation Loss: 4.135494	Validation
Epoch: 168	Training Loss: 3.816918	Validation Loss: 4.081820	Validation
Epoch: 169	Training Loss: 3.828218	Validation Loss: 4.104443	Validation
Epoch: 170	Training Loss: 3.802220	Validation Loss: 4.152456	Validation
Epoch: 171	Training Loss: 3.829246	Validation Loss: 4.092724	Validation
Epoch: 172	Training Loss: 3.786911	Validation Loss: 4.095738	Validation
Epoch: 173	Training Loss: 3.784518	Validation Loss: 4.118988	Validation
Epoch: 174	Training Loss: 3.771987	Validation Loss: 4.116381	Validation
Epoch: 175	Training Loss: 3.765835	Validation Loss: 4.147020	Validation
Epoch: 176	Training Loss: 3.770134	Validation Loss: 4.106996	Validation
Epoch: 177	Training Loss: 3.770224	Validation Loss: 4.116591	Validation
Epoch: 178	Training Loss: 3.728404	Validation Loss: 4.077125	Validation
Epoch: 179	Training Loss: 3.741863	Validation Loss: 4.124872	Validation
Epoch: 180	Training Loss: 3.759200	Validation Loss: 4.086443	Validation
Epoch: 181	Training Loss: 3.728881	Validation Loss: 4.139641	Validation
Epoch: 182	Training Loss: 3.719232	Validation Loss: 4.071069	Validation
Epoch: 183	Training Loss: 3.735832	Validation Loss: 4.134961	Validation
Epoch: 184	Training Loss: 3.694988	Validation Loss: 4.057529	Validation
Epoch: 185	Training Loss: 3.711285	Validation Loss: 4.045784	Validation
Epoch: 186	Training Loss: 3.700440	Validation Loss: 4.035214	Validation
Epoch: 187	Training Loss: 3.713903	Validation Loss: 4.069977	Validation
Epoch: 188	Training Loss: 3.708422	Validation Loss: 4.087810	Validation
Epoch: 189	Training Loss: 3.682284	Validation Loss: 4.072077	Validation
Epoch: 190	Training Loss: 3.701727	Validation Loss: 4.060647	Validation
Epoch: 191	Training Loss: 3.692966	Validation Loss: 4.067240	Validation
Epoch: 192	Training Loss: 3.676061	Validation Loss: 4.039976	Validation
Epoch: 193	Training Loss: 3.675046	Validation Loss: 4.082298	Validation
Epoch: 194	Training Loss: 3.656907	Validation Loss: 4.047762	Validation
Epoch: 195	Training Loss: 3.659859	Validation Loss: 4.073039	Validation
Epoch: 196	Training Loss: 3.669619	Validation Loss: 4.030258	Validation
Epoch: 197	Training Loss: 3.636013	Validation Loss: 4.021617	Validation
Epoch: 198	Training Loss: 3.607924	Validation Loss: 4.021764	Validation
Epoch: 199	Training Loss: 3.653014	Validation Loss: 4.027783	Validation
Epoch: 200	Training Loss: 3.598336	Validation Loss: 4.039530	Validation
Epoch: 201	Training Loss: 3.615581	Validation Loss: 4.020327	Validation
Epoch: 202	Training Loss: 3.592008	Validation Loss: 4.014199	Validation
Epoch: 203	Training Loss: 3.584552	Validation Loss: 4.053055	Validation
Epoch: 204	Training Loss: 3.605174	Validation Loss: 4.068951	Validation
Epoch: 205	Training Loss: 3.576563	Validation Loss: 3.976563	Validation
Epoch: 206	Training Loss: 3.595612	Validation Loss: 4.046586	Validation
Epoch: 207	Training Loss: 3.583784	Validation Loss: 3.976213	Validation
Epoch: 208	Training Loss: 3.582941	Validation Loss: 4.012356	Validation
Epoch: 209	Training Loss: 3.571632	Validation Loss: 3.970572	Validation
Epoch: 210	Training Loss: 3.559407	Validation Loss: 4.058072	Validation
Epoch: 211	Training Loss: 3.592970	Validation Loss: 4.000300	Validation

Epoch: 212	Training Loss: 3.589015	Validation Loss: 4.030813	Validation
Epoch: 213	Training Loss: 3.547368	Validation Loss: 4.051563	Validation
Epoch: 214	Training Loss: 3.559373	Validation Loss: 4.007505	Validation
Epoch: 215	Training Loss: 3.538662	Validation Loss: 3.983643	Validation
Epoch: 216	Training Loss: 3.528767	Validation Loss: 4.038404	Validation
Epoch: 217	Training Loss: 3.517386	Validation Loss: 3.977580	Validation
Epoch: 218	Training Loss: 3.531313	Validation Loss: 4.012578	Validation
Epoch: 219	Training Loss: 3.520049	Validation Loss: 4.008932	Validation
Epoch: 220	Training Loss: 3.495922	Validation Loss: 3.974050	Validation
Epoch: 221	Training Loss: 3.479124	Validation Loss: 4.022565	Validation
Epoch: 222	Training Loss: 3.506943	Validation Loss: 3.999432	Validation
Epoch: 223	Training Loss: 3.496641	Validation Loss: 3.949164	Validation
Epoch: 224	Training Loss: 3.492117	Validation Loss: 4.068584	Validation
Epoch: 225	Training Loss: 3.479025	Validation Loss: 4.010494	Validation
Epoch: 226	Training Loss: 3.479768	Validation Loss: 4.038422	Validation
Epoch: 227	Training Loss: 3.473467	Validation Loss: 4.005161	Validation
Epoch: 228	Training Loss: 3.475405	Validation Loss: 4.008047	Validation
Epoch: 229	Training Loss: 3.431635	Validation Loss: 4.037668	Validation
Epoch: 230	Training Loss: 3.432604	Validation Loss: 3.944960	Validation
Epoch: 231	Training Loss: 3.478246	Validation Loss: 4.001677	Validation
Epoch: 232	Training Loss: 3.448058	Validation Loss: 3.993002	Validation
Epoch: 233	Training Loss: 3.414895	Validation Loss: 3.961176	Validation
Epoch: 234	Training Loss: 3.423721	Validation Loss: 4.006955	Validation
Epoch: 235	Training Loss: 3.439458	Validation Loss: 3.960396	Validation
Epoch: 236	Training Loss: 3.393210	Validation Loss: 3.991027	Validation
Epoch: 237	Training Loss: 3.429129	Validation Loss: 3.977577	Validation
Epoch: 238	Training Loss: 3.393116	Validation Loss: 3.981567	Validation
Epoch: 239	Training Loss: 3.419396	Validation Loss: 4.003803	Validation
Epoch: 240	Training Loss: 3.413674	Validation Loss: 3.884477	Validation
Epoch: 241	Training Loss: 3.409231	Validation Loss: 4.008877	Validation
Epoch: 242	Training Loss: 3.380367	Validation Loss: 3.960272	Validation
Epoch: 243	Training Loss: 3.377679	Validation Loss: 3.981557	Validation
Epoch: 244	Training Loss: 3.351337	Validation Loss: 3.927392	Validation
Epoch: 245	Training Loss: 3.337668	Validation Loss: 4.004786	Validation
Epoch: 246	Training Loss: 3.342791	Validation Loss: 4.017981	Validation
Epoch: 247	Training Loss: 3.344606	Validation Loss: 3.989302	Validation
Epoch: 248	Training Loss: 3.356205	Validation Loss: 3.940933	Validation
Epoch: 249	Training Loss: 3.356083	Validation Loss: 3.938275	Validation
Epoch: 250	Training Loss: 3.333505	Validation Loss: 3.993502	Validation
Epoch: 251	Training Loss: 3.325341	Validation Loss: 4.013399	Validation
Epoch: 252	Training Loss: 3.319654	Validation Loss: 3.938750	Validation
Epoch: 253	Training Loss: 3.300279	Validation Loss: 3.995460	Validation
Epoch: 254	Training Loss: 3.315623	Validation Loss: 4.005507	Validation
Epoch: 255	Training Loss: 3.316379	Validation Loss: 3.943687	Validation
Epoch: 256	Training Loss: 3.312566	Validation Loss: 3.950633	Validation
Epoch: 257	Training Loss: 3.301498	Validation Loss: 4.038296	Validation
Epoch: 258	Training Loss: 3.296878	Validation Loss: 3.912035	Validation
Epoch: 259	Training Loss: 3.292913	Validation Loss: 4.019599	Validation

Epoch: 260	Training Loss: 3.299337	Validation Loss: 3.950137	Validation
Epoch: 261	Training Loss: 3.292271	Validation Loss: 3.933327	Validation

1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```
In [0]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['test']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
        100. * correct / total, correct, total))

    # call test function
    test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)
```

Test Loss: 3.916158

Test Accuracy: 11% (100/836)

Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```
In [0]: ## TODO: Specify data loaders
import os
from torchvision import datasets
from torchvision import transforms
from torch.utils import data

import matplotlib.pyplot as plt

data_transform = transforms.Compose([
    transforms.RandomResizedCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

data_transform_train = transforms.Compose([
    transforms.RandomResizedCrop(224),
    transforms.RandomRotation((0, 90)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

loaders_transfer = {}
data_transfer = {}
batch_size = 10
num_workers = 2

train_dataset = datasets.ImageFolder(root='/content/drive/My Drive/Colab Notebooks/data/dogImages/train')
data_transfer['train'] = train_dataset
loaders_transfer['train'] = data.DataLoader(train_dataset, batch_size=batch_size, shuffle=True)

valid_dataset = datasets.ImageFolder(root='/content/drive/My Drive/Colab Notebooks/data/dogImages/valid')
data_transfer['valid'] = valid_dataset
loaders_transfer['valid'] = data.DataLoader(valid_dataset, batch_size=batch_size, shuffle=True)

test_dataset = datasets.ImageFolder(root='/content/drive/My Drive/Colab Notebooks/data/dogImages/test')
data_transfer['test'] = test_dataset
```

```
loaders_transfer['test'] = data.DataLoader(test_dataset, batch_size=batch_size, shuffle=
```

1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```
In [0]: import torchvision.models as models
import torch.nn as nn

## TODO: Specify model architecture
model_transfer = models.vgg16(pretrained=True)

for param in model_transfer.features.parameters():
    param.requires_grad = False

model_transfer.classifier[6] = nn.Linear(model_transfer.classifier[6].in_features, 133)

if use_cuda:
    model_transfer = model_transfer.cuda()
```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer: I looked on models available in PyTorch and decided to use VGG16 model. I selected that model because it is suitable for image recognition and among other things it recognizes dog's breeds as well. Therefore set of features it calculates should contain some dog breeds characteristics. I decided to modify only last fully connected layer, so it returns 133 classes, which correspond to dog's breeds. Then I initialized learning process only for fully connected layer.

1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [0]: import torch.optim as optim

learning_rate = 0.001

### TODO: select loss function
criterion_transfer = nn.CrossEntropyLoss()

### TODO: select optimizer
optimizer_transfer = optim.SGD(model_transfer.parameters(), lr=learning_rate)
```

1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```

In [0]: # the following import is required for training to be robust to truncated images
import numpy as np
from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True

def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf
    correct = 0.
    total = 0.

    for epoch in range(1, n_epochs+1):
        # initialize variables to monitor training and validation loss
        train_loss = 0.0
        valid_loss = 0.0

        #####
        # train the model #
        #####
        model.train()
        for batch_idx, (data, target) in enumerate(loaders['train']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            ## find the loss and update the model parameters accordingly
            ## record the average training loss, using something like
            ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

            # clear the gradients of all optimized variables
            optimizer.zero_grad()
            # forward pass: compute predicted outputs by passing inputs to the model
            output = model(data)
            # calculate the batch loss
            loss = criterion(output, target)
            # backward pass: compute gradient of the loss with respect to model parameters
            loss.backward()
            # perform a single optimization step (parameter update)
            optimizer.step()
            # update training loss
            train_loss += ((1 / (batch_idx + 1)) * (loss.data - train_loss))

        #####
        # validate the model #
        #####
        model.eval()
        for batch_idx, (data, target) in enumerate(loaders['valid']):

```

```

# move to GPU
if use_cuda:
    data, target = data.cuda(), target.cuda()
## update the average validation loss

# forward pass: compute predicted outputs by passing inputs to the model
output = model(data)
# calculate the batch loss
loss = criterion(output, target)
# update average validation loss
valid_loss += ((1 / (batch_idx + 1)) * (loss.data - valid_loss))
# convert output probabilities to predicted class
pred = output.data.max(1, keepdim=True)[1]
# compare predictions to true label
correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy)
total += data.size(0)

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f} \tValidation
      epoch,
      train_loss,
      valid_loss,
      100. * correct / total
    ))

## TODO: save the model if validation loss has decreased
if valid_loss < valid_loss_min:
    torch.save(model.state_dict(), save_path)
    valid_loss_min = valid_loss

# return trained model
return model

# train the model
model_transfer = train(200, loaders_transfer, model_transfer, optimizer_transfer, criterion_transfer)

# load the model that got the best validation accuracy (uncomment the line below)
model_transfer.load_state_dict(torch.load('/content/drive/My Drive/Colab Notebooks/model_transfer.pth'))

```

Epoch: 1	Training Loss: 4.350549	Validation Loss: 3.048523	Validation Accuracy: 0.000000
Epoch: 2	Training Loss: 3.022008	Validation Loss: 1.852071	Validation Accuracy: 0.000000
Epoch: 3	Training Loss: 2.436795	Validation Loss: 1.395277	Validation Accuracy: 0.000000
Epoch: 4	Training Loss: 2.190438	Validation Loss: 1.250793	Validation Accuracy: 0.000000
Epoch: 5	Training Loss: 2.009026	Validation Loss: 1.243276	Validation Accuracy: 0.000000
Epoch: 6	Training Loss: 1.988504	Validation Loss: 1.092259	Validation Accuracy: 0.000000
Epoch: 7	Training Loss: 1.854950	Validation Loss: 1.067003	Validation Accuracy: 0.000000
Epoch: 8	Training Loss: 1.820828	Validation Loss: 1.048521	Validation Accuracy: 0.000000

Epoch: 9	Training Loss: 1.793996	Validation Loss: 0.986793	Validation Accuracy: 0.811483
Epoch: 10	Training Loss: 1.756487	Validation Loss: 1.014400	Validation Accuracy: 0.811483
Epoch: 11	Training Loss: 1.709455	Validation Loss: 0.993467	Validation Accuracy: 0.811483
Epoch: 12	Training Loss: 1.725079	Validation Loss: 0.972734	Validation Accuracy: 0.811483
Epoch: 13	Training Loss: 1.709674	Validation Loss: 0.978867	Validation Accuracy: 0.811483
Epoch: 14	Training Loss: 1.660839	Validation Loss: 0.950839	Validation Accuracy: 0.811483
Epoch: 15	Training Loss: 1.631993	Validation Loss: 0.916922	Validation Accuracy: 0.811483
Epoch: 16	Training Loss: 1.590955	Validation Loss: 0.944802	Validation Accuracy: 0.811483
Epoch: 17	Training Loss: 1.574732	Validation Loss: 0.954425	Validation Accuracy: 0.811483
Epoch: 18	Training Loss: 1.574804	Validation Loss: 1.034662	Validation Accuracy: 0.811483
Epoch: 19	Training Loss: 1.569586	Validation Loss: 0.959205	Validation Accuracy: 0.811483
Epoch: 20	Training Loss: 1.528988	Validation Loss: 0.866728	Validation Accuracy: 0.811483
Epoch: 21	Training Loss: 1.524658	Validation Loss: 0.936819	Validation Accuracy: 0.811483
Epoch: 22	Training Loss: 1.513015	Validation Loss: 0.875037	Validation Accuracy: 0.811483
Epoch: 23	Training Loss: 1.481418	Validation Loss: 0.883922	Validation Accuracy: 0.811483
Epoch: 24	Training Loss: 1.487957	Validation Loss: 0.950757	Validation Accuracy: 0.811483
Epoch: 25	Training Loss: 1.473754	Validation Loss: 0.922161	Validation Accuracy: 0.811483
Epoch: 26	Training Loss: 1.455523	Validation Loss: 0.979282	Validation Accuracy: 0.811483
Epoch: 27	Training Loss: 1.466012	Validation Loss: 0.876578	Validation Accuracy: 0.811483
Epoch: 28	Training Loss: 1.459650	Validation Loss: 0.919370	Validation Accuracy: 0.811483
Epoch: 29	Training Loss: 1.423060	Validation Loss: 0.860250	Validation Accuracy: 0.811483
Epoch: 30	Training Loss: 1.405744	Validation Loss: 0.899574	Validation Accuracy: 0.811483
Epoch: 31	Training Loss: 1.432563	Validation Loss: 0.856633	Validation Accuracy: 0.811483
Epoch: 32	Training Loss: 1.400478	Validation Loss: 0.878073	Validation Accuracy: 0.811483
Epoch: 33	Training Loss: 1.383176	Validation Loss: 0.868487	Validation Accuracy: 0.811483
Epoch: 34	Training Loss: 1.372401	Validation Loss: 0.911573	Validation Accuracy: 0.811483
Epoch: 35	Training Loss: 1.379999	Validation Loss: 0.862773	Validation Accuracy: 0.811483
Epoch: 36	Training Loss: 1.382046	Validation Loss: 0.867083	Validation Accuracy: 0.811483
Epoch: 37	Training Loss: 1.365462	Validation Loss: 0.864638	Validation Accuracy: 0.811483
Epoch: 38	Training Loss: 1.363595	Validation Loss: 0.969619	Validation Accuracy: 0.811483
Epoch: 39	Training Loss: 1.348352	Validation Loss: 0.888406	Validation Accuracy: 0.811483
Epoch: 40	Training Loss: 1.358488	Validation Loss: 0.828097	Validation Accuracy: 0.811483
Epoch: 41	Training Loss: 1.341664	Validation Loss: 0.826655	Validation Accuracy: 0.811483
Epoch: 42	Training Loss: 1.343980	Validation Loss: 0.858087	Validation Accuracy: 0.811483
Epoch: 43	Training Loss: 1.320033	Validation Loss: 0.906909	Validation Accuracy: 0.811483
Epoch: 44	Training Loss: 1.342481	Validation Loss: 0.932043	Validation Accuracy: 0.811483
Epoch: 45	Training Loss: 1.288690	Validation Loss: 0.900238	Validation Accuracy: 0.811483
Epoch: 46	Training Loss: 1.321352	Validation Loss: 0.915504	Validation Accuracy: 0.811483
Epoch: 47	Training Loss: 1.320359	Validation Loss: 0.858901	Validation Accuracy: 0.811483
Epoch: 48	Training Loss: 1.293133	Validation Loss: 0.871253	Validation Accuracy: 0.811483
Epoch: 49	Training Loss: 1.280491	Validation Loss: 0.811483	Validation Accuracy: 0.811483
Epoch: 50	Training Loss: 1.277678	Validation Loss: 0.842175	Validation Accuracy: 0.811483
Epoch: 51	Training Loss: 1.272234	Validation Loss: 0.918432	Validation Accuracy: 0.811483
Epoch: 52	Training Loss: 1.247608	Validation Loss: 0.861602	Validation Accuracy: 0.811483
Epoch: 53	Training Loss: 1.294649	Validation Loss: 0.875606	Validation Accuracy: 0.811483
Epoch: 54	Training Loss: 1.243491	Validation Loss: 0.786732	Validation Accuracy: 0.811483
Epoch: 55	Training Loss: 1.243947	Validation Loss: 0.847915	Validation Accuracy: 0.811483
Epoch: 56	Training Loss: 1.225638	Validation Loss: 0.910021	Validation Accuracy: 0.811483

Epoch: 57	Training Loss: 1.229640	Validation Loss: 0.833956	Validation A
Epoch: 58	Training Loss: 1.246513	Validation Loss: 0.822229	Validation A
Epoch: 59	Training Loss: 1.233181	Validation Loss: 0.874539	Validation A
Epoch: 60	Training Loss: 1.251836	Validation Loss: 0.839039	Validation A
Epoch: 61	Training Loss: 1.228976	Validation Loss: 0.916061	Validation A
Epoch: 62	Training Loss: 1.205238	Validation Loss: 0.814595	Validation A
Epoch: 63	Training Loss: 1.219204	Validation Loss: 0.877125	Validation A
Epoch: 64	Training Loss: 1.184932	Validation Loss: 0.833166	Validation A
Epoch: 65	Training Loss: 1.179863	Validation Loss: 0.882737	Validation A
Epoch: 66	Training Loss: 1.178546	Validation Loss: 0.894260	Validation A
Epoch: 67	Training Loss: 1.179698	Validation Loss: 0.821752	Validation A
Epoch: 68	Training Loss: 1.175331	Validation Loss: 0.906829	Validation A
Epoch: 69	Training Loss: 1.188597	Validation Loss: 0.785788	Validation A
Epoch: 70	Training Loss: 1.166398	Validation Loss: 0.923929	Validation A
Epoch: 71	Training Loss: 1.155743	Validation Loss: 0.892187	Validation A
Epoch: 72	Training Loss: 1.179028	Validation Loss: 0.877222	Validation A
Epoch: 73	Training Loss: 1.167311	Validation Loss: 0.913024	Validation A
Epoch: 74	Training Loss: 1.169264	Validation Loss: 0.902663	Validation A
Epoch: 75	Training Loss: 1.161647	Validation Loss: 0.804652	Validation A
Epoch: 76	Training Loss: 1.155036	Validation Loss: 0.885070	Validation A
Epoch: 77	Training Loss: 1.164256	Validation Loss: 0.869039	Validation A
Epoch: 78	Training Loss: 1.151888	Validation Loss: 0.904351	Validation A
Epoch: 79	Training Loss: 1.120381	Validation Loss: 0.879150	Validation A
Epoch: 80	Training Loss: 1.137585	Validation Loss: 0.887680	Validation A
Epoch: 81	Training Loss: 1.133347	Validation Loss: 0.887100	Validation A
Epoch: 82	Training Loss: 1.124412	Validation Loss: 0.897556	Validation A
Epoch: 83	Training Loss: 1.128219	Validation Loss: 0.837869	Validation A
Epoch: 84	Training Loss: 1.129614	Validation Loss: 0.828445	Validation A
Epoch: 85	Training Loss: 1.113989	Validation Loss: 0.865107	Validation A
Epoch: 86	Training Loss: 1.141778	Validation Loss: 0.871239	Validation A
Epoch: 87	Training Loss: 1.121677	Validation Loss: 0.856440	Validation A
Epoch: 88	Training Loss: 1.124460	Validation Loss: 0.818129	Validation A
Epoch: 89	Training Loss: 1.119911	Validation Loss: 0.821620	Validation A
Epoch: 90	Training Loss: 1.127928	Validation Loss: 0.914855	Validation A
Epoch: 91	Training Loss: 1.074187	Validation Loss: 0.828309	Validation A
Epoch: 92	Training Loss: 1.111570	Validation Loss: 0.864336	Validation A
Epoch: 93	Training Loss: 1.090127	Validation Loss: 0.852935	Validation A
Epoch: 94	Training Loss: 1.093921	Validation Loss: 0.890560	Validation A
Epoch: 95	Training Loss: 1.114071	Validation Loss: 0.818762	Validation A
Epoch: 96	Training Loss: 1.110500	Validation Loss: 0.871621	Validation A
Epoch: 97	Training Loss: 1.071164	Validation Loss: 0.906091	Validation A
Epoch: 98	Training Loss: 1.078476	Validation Loss: 0.847184	Validation A
Epoch: 99	Training Loss: 1.090338	Validation Loss: 0.873513	Validation A
Epoch: 100	Training Loss: 1.096459	Validation Loss: 0.825962	Validation A
Epoch: 101	Training Loss: 1.051832	Validation Loss: 0.798784	Validation A
Epoch: 102	Training Loss: 1.058251	Validation Loss: 0.874341	Validation A
Epoch: 103	Training Loss: 1.065672	Validation Loss: 0.838769	Validation A
Epoch: 104	Training Loss: 1.067207	Validation Loss: 0.787508	Validation A

Epoch: 105	Training Loss: 1.069818	Validation Loss: 0.852506	Validation
Epoch: 106	Training Loss: 1.033501	Validation Loss: 0.891365	Validation
Epoch: 107	Training Loss: 1.056100	Validation Loss: 0.869576	Validation
Epoch: 108	Training Loss: 1.053541	Validation Loss: 0.835647	Validation
Epoch: 109	Training Loss: 1.058026	Validation Loss: 0.851428	Validation
Epoch: 110	Training Loss: 1.049820	Validation Loss: 0.872159	Validation
Epoch: 111	Training Loss: 1.029159	Validation Loss: 0.801336	Validation
Epoch: 112	Training Loss: 1.052759	Validation Loss: 0.829640	Validation
Epoch: 113	Training Loss: 1.048656	Validation Loss: 0.884582	Validation
Epoch: 114	Training Loss: 1.037304	Validation Loss: 0.784785	Validation
Epoch: 115	Training Loss: 1.029194	Validation Loss: 0.825595	Validation
Epoch: 116	Training Loss: 1.021951	Validation Loss: 0.835324	Validation
Epoch: 117	Training Loss: 1.037267	Validation Loss: 0.909042	Validation
Epoch: 118	Training Loss: 1.034168	Validation Loss: 0.842677	Validation
Epoch: 119	Training Loss: 1.011630	Validation Loss: 0.863215	Validation
Epoch: 120	Training Loss: 1.053097	Validation Loss: 0.802961	Validation
Epoch: 121	Training Loss: 1.020529	Validation Loss: 0.907416	Validation
Epoch: 122	Training Loss: 1.041613	Validation Loss: 0.964690	Validation
Epoch: 123	Training Loss: 1.024317	Validation Loss: 0.855272	Validation
Epoch: 124	Training Loss: 1.022575	Validation Loss: 0.839153	Validation
Epoch: 125	Training Loss: 1.032861	Validation Loss: 0.913314	Validation
Epoch: 126	Training Loss: 0.984614	Validation Loss: 0.869581	Validation
Epoch: 127	Training Loss: 1.035380	Validation Loss: 0.849176	Validation
Epoch: 128	Training Loss: 0.980110	Validation Loss: 0.872390	Validation
Epoch: 129	Training Loss: 0.976327	Validation Loss: 0.876102	Validation
Epoch: 130	Training Loss: 0.987447	Validation Loss: 0.766225	Validation
Epoch: 131	Training Loss: 1.022517	Validation Loss: 0.827656	Validation
Epoch: 132	Training Loss: 0.971149	Validation Loss: 0.798763	Validation
Epoch: 133	Training Loss: 1.009829	Validation Loss: 0.876827	Validation
Epoch: 134	Training Loss: 0.960093	Validation Loss: 0.937074	Validation
Epoch: 135	Training Loss: 0.977952	Validation Loss: 0.805697	Validation
Epoch: 136	Training Loss: 0.945256	Validation Loss: 0.848456	Validation
Epoch: 137	Training Loss: 0.984060	Validation Loss: 0.827487	Validation
Epoch: 138	Training Loss: 0.967931	Validation Loss: 0.863880	Validation
Epoch: 139	Training Loss: 0.993942	Validation Loss: 0.838410	Validation
Epoch: 140	Training Loss: 0.958190	Validation Loss: 0.831148	Validation
Epoch: 141	Training Loss: 0.958597	Validation Loss: 0.880443	Validation
Epoch: 142	Training Loss: 0.958433	Validation Loss: 0.770863	Validation
Epoch: 143	Training Loss: 0.945799	Validation Loss: 0.839779	Validation
Epoch: 144	Training Loss: 0.973686	Validation Loss: 0.836786	Validation
Epoch: 145	Training Loss: 0.967472	Validation Loss: 0.922112	Validation
Epoch: 146	Training Loss: 0.984392	Validation Loss: 0.864949	Validation
Epoch: 147	Training Loss: 0.984810	Validation Loss: 0.745232	Validation
Epoch: 148	Training Loss: 0.966439	Validation Loss: 0.841016	Validation
Epoch: 149	Training Loss: 0.942397	Validation Loss: 0.824959	Validation
Epoch: 150	Training Loss: 0.966647	Validation Loss: 0.848369	Validation
Epoch: 151	Training Loss: 0.957443	Validation Loss: 0.834324	Validation
Epoch: 152	Training Loss: 0.959414	Validation Loss: 0.818835	Validation

Epoch: 153	Training Loss: 0.927498	Validation Loss: 0.851661	Validation
Epoch: 154	Training Loss: 0.979422	Validation Loss: 0.913471	Validation
Epoch: 155	Training Loss: 0.962420	Validation Loss: 0.890850	Validation
Epoch: 156	Training Loss: 0.925692	Validation Loss: 0.879071	Validation
Epoch: 157	Training Loss: 0.965130	Validation Loss: 0.847183	Validation
Epoch: 158	Training Loss: 0.950044	Validation Loss: 0.808831	Validation
Epoch: 159	Training Loss: 0.943548	Validation Loss: 0.830338	Validation
Epoch: 160	Training Loss: 0.905804	Validation Loss: 0.882916	Validation
Epoch: 161	Training Loss: 0.937819	Validation Loss: 0.866552	Validation
Epoch: 162	Training Loss: 0.894047	Validation Loss: 0.833897	Validation
Epoch: 163	Training Loss: 0.926091	Validation Loss: 0.858849	Validation
Epoch: 164	Training Loss: 0.944304	Validation Loss: 0.800117	Validation
Epoch: 165	Training Loss: 0.945319	Validation Loss: 0.855845	Validation
Epoch: 166	Training Loss: 0.943876	Validation Loss: 0.896114	Validation
Epoch: 167	Training Loss: 0.919474	Validation Loss: 0.840434	Validation
Epoch: 168	Training Loss: 0.948620	Validation Loss: 0.936812	Validation
Epoch: 169	Training Loss: 0.930282	Validation Loss: 0.911792	Validation
Epoch: 170	Training Loss: 0.910861	Validation Loss: 0.853065	Validation
Epoch: 171	Training Loss: 0.947704	Validation Loss: 0.871225	Validation
Epoch: 172	Training Loss: 0.912299	Validation Loss: 0.870394	Validation
Epoch: 173	Training Loss: 0.927816	Validation Loss: 0.837104	Validation
Epoch: 174	Training Loss: 0.922887	Validation Loss: 0.953049	Validation
Epoch: 175	Training Loss: 0.917136	Validation Loss: 0.885836	Validation
Epoch: 176	Training Loss: 0.933738	Validation Loss: 0.881380	Validation
Epoch: 177	Training Loss: 0.895659	Validation Loss: 0.781877	Validation
Epoch: 178	Training Loss: 0.901494	Validation Loss: 0.782590	Validation
Epoch: 179	Training Loss: 0.959918	Validation Loss: 0.861145	Validation
Epoch: 180	Training Loss: 0.881766	Validation Loss: 0.875878	Validation
Epoch: 181	Training Loss: 0.879401	Validation Loss: 0.841445	Validation
Epoch: 182	Training Loss: 0.895581	Validation Loss: 0.872405	Validation
Epoch: 183	Training Loss: 0.922558	Validation Loss: 0.861642	Validation
Epoch: 184	Training Loss: 0.901537	Validation Loss: 0.935880	Validation
Epoch: 185	Training Loss: 0.869326	Validation Loss: 0.809489	Validation
Epoch: 186	Training Loss: 0.865229	Validation Loss: 0.878388	Validation
Epoch: 187	Training Loss: 0.912063	Validation Loss: 0.826799	Validation
Epoch: 188	Training Loss: 0.904545	Validation Loss: 0.866770	Validation
Epoch: 189	Training Loss: 0.892413	Validation Loss: 0.809744	Validation
Epoch: 190	Training Loss: 0.885405	Validation Loss: 0.875728	Validation
Epoch: 191	Training Loss: 0.885403	Validation Loss: 0.869741	Validation
Epoch: 192	Training Loss: 0.899547	Validation Loss: 0.801121	Validation
Epoch: 193	Training Loss: 0.871854	Validation Loss: 0.894483	Validation
Epoch: 194	Training Loss: 0.882282	Validation Loss: 0.823136	Validation
Epoch: 195	Training Loss: 0.885834	Validation Loss: 0.916837	Validation
Epoch: 196	Training Loss: 0.886272	Validation Loss: 0.815705	Validation
Epoch: 197	Training Loss: 0.829622	Validation Loss: 0.820222	Validation
Epoch: 198	Training Loss: 0.874619	Validation Loss: 0.852219	Validation
Epoch: 199	Training Loss: 0.871274	Validation Loss: 0.875762	Validation
Epoch: 200	Training Loss: 0.882420	Validation Loss: 0.890313	Validation

```
Out[0]: <All keys matched successfully>
```

1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [0]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['test']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
        100. * correct / total, correct, total))

    test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

```
Test Loss: 0.835999
```

```
Test Accuracy: 77% (650/836)
```

1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```

In [0]: ### TODO: Write a function that takes a path to an image as input
        ### and returns the dog breed that is predicted by the model.

        from PIL import Image
        import torchvision.transforms as transforms
        from torch.autograd import Variable

        # Set PIL to be tolerant of image files that are truncated.
        from PIL import ImageFile
        ImageFile.LOAD_TRUNCATED_IMAGES = True

        # list of class names by index, i.e. a name can be accessed like class_names[0]
        class_names = [item[4:].replace("_", " ") for item in data_transfer['train'].classes]

        def predict_breed_transfer(img_path):
            # load the image and return the predicted breed

            image = Image.open(img_path)
            transform = transforms.Compose([
                transforms.ToTensor(),
                transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
            ])

            image = transform(image)
            image = image.unsqueeze(0)
            if use_cuda:
                image = image.cuda()
            prediction = model_transfer(image)
            prediction_class = np.argmax(prediction.cpu().detach().numpy()) # predicted class in

            return class_names[prediction_class]

```

Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `dog_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!



Sample Human Output

1.1.18 (IMPLEMENTATION) Write your Algorithm

```
In [0]: ### TODO: Write your algorithm.
        ### Feel free to use as many code cells as needed.

def run_app(img_path):
    ## handle cases for a human face, dog, and neither

    isFace = face_detector(img_path)
    isDog = dog_detector(img_path)
    if isDog:
        breed = predict_breed_transfer(img_path)
        return "It is a dog. Detected breed is: " + breed
    elif isFace:
        breed = predict_breed_transfer(img_path)
        return "It is a human. Detected breed is: " + breed
    else:
        return "It is neither dog nor human."
```

Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: I tested network on some cats, dogs, and human images. The network did not detect cats neither as dogs nor as human, which is very good result. The most of dog breeds were recognized correctly. On the other hand humans were not really similar to detected dog's breed.

To improve algorithm we could try the following: - increase number of samples used for learning the network - increase number of features (convolutional layer) detected by the network - add additional fully connected layer - use technologies of model ensembling

```
In [3]: import numpy as np
        from glob import glob

        # load filenames for human and dog images
        cat_files = np.array(glob("/content/drive/My Drive/Colab Notebooks/data/cats/*"))
        human_files = np.array(glob("/content/drive/My Drive/Colab Notebooks/data/humans/*"))
        dog_files = np.array(glob("/content/drive/My Drive/Colab Notebooks/data/dogs/*"))

        # print number of images in each dataset
        print('There are %d total cat images.' % len(cat_files))
        print('There are %d total human images.' % len(human_files))
        print('There are %d total dog images.' % len(dog_files))
```


There are 4 total cat images.

There are 4 total human images.


There are 4 total dog images.

```
In [22]: ## TODO: Execute your algorithm from Step 6 on
        ## at least 6 images on your computer.
        ## Feel free to use as many code cells as needed.
        import cv2
        import matplotlib.pyplot as plt
        %matplotlib inline

        ## suggested code, below
        for file in np.hstack((cat_files, human_files, dog_files)):
            img = cv2.imread(file)
            cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
            plt.imshow(cv_rgb)
            plt.show()
            result = run_app(file)
            print(result)
```



It is neither dog nor human.



It is neither dog nor human.



It is neither dog nor human.




It is neither dog nor human.



It is a human. Detected breed is: Akita




It is a human. Detected breed is: Akita




It is a human. Detected breed is: Chinese shar-pei



It is a human. Detected breed is: Chinese crested



It is a dog. Detected breed is: Great dane



It is a dog. Detected breed is: Anatolian shepherd dog



It is a dog. Detected breed is: Otterhound



It is a dog. Detected breed is: German shepherd dog